

## Assignment-3

### CS540- Database Management Systems

Arash Termehchy

**Source code:** WordCount.java, Map.Java, Reduce.Java

Wordcount.java:

```
package kollipal_Wordcount;

import java.util.*;
import java.io.*;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;

public class WordCount {
    public static void main(String[] args) throws Exception
    {
        Configuration conf = new Configuration();
        Job job = new Job(conf, "kollipal_WordCount");
        job.setJarByClass(WordCount.class);
        //Job job = Job.getInstance(conf);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setJarByClass(Map.class);
        job.setMapperClass(Map.class);
        job.setJarByClass(Reduce.class);
        job.setReducerClass(Reduce.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.waitForCompletion(true);
    }
}
```

Lakshman Madhav Kollipara  
ONID-kollipal  
CS540 – Database Management Systems  
Map.Java:

```
package kollipal_Wordcount;

import java.util.*;
import java.io.*;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class Map extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void map(LongWritable key, Text value, Context context)
    throws IOException, InterruptedException
    {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens())
        {
            word.set(tokenizer.nextToken());
            context.write(word, one);
        }
    }
}
```

Lakshman Madhav Kollipara  
ONID-kollipal  
CS540 – Database Management Systems  
Reduce.java:

```
package kollipal_Wordcount;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.Reducer.Context;
import java.io.*;
import java.util.*;

public class Reduce extends Reducer<Text,IntWritable,Text,IntWritable>
{
    int max_sum = 0;
    Text max_occured_key = new Text();
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException
    {
        int sum = 0;

        for(IntWritable value: values){
            sum+= value.get();
        }
        if (sum > max_sum)
        {
            max_sum = sum;
            max_occured_key.set(key);
        }
        //context.write(key, new IntWritable(sum));
    }

    @Override
    protected void cleanup(Context context)
        throws IOException,InterruptedException
    {
        context.write(max_occured_key, new
IntWritable(max_sum));
    }
}
```

**JAR file:** kollipal\_WordCount.Jar

**Hadoop Version:** Hadoop 2.6.0

**HDP Version:** HDP 2.2

**Input:**

File-1.txt

Lakshman  
Madhav  
Kollipara  
Oregon  
State  
University  
Corvallis  
Lakshman  
Madhav  
Kollipara  
Oregon  
State  
University  
Corvallis  
Lakshman  
Lakshman

File-2.txt

Lakshman  
Madhav  
Kollipara  
Oregon  
State  
University  
Corvallis  
Lakshman  
Madhav  
Kollipara  
Oregon  
State  
University  
Corvallis  
Lakshman

**Output:**

Lakshman      7

Intermediate Step: (Total occurrence of each word)

Corvallis      4  
Kollipara      4  
Lakshman      7  
Madhav      4  
Oregon      4  
State      4  
University      4

The screenshot shows the HDFS web interface. At the top, the breadcrumb path is 'Home / user / hue / Output\_1 / part-r-00000'. On the left, there is a sidebar with 'ACTIONS' (View As Binary, Edit File, Download, View File Location, Refresh) and 'INFO'. The main area displays file details: 'First Block', 'Previous Block', 'Next Block', 'Last Block', and 'Viewing Bytes: 1 - 11 of 11 (4096 B block size)'. Below this, a table shows the file 'Lakshman' with a size of 7. At the bottom, there are more navigation buttons: 'First Block', 'Previous Block', 'Next Block', 'Last Block'.

## Job Browser: 100% map and 100% reduce Successful

The screenshot shows the Hadoop Job Browser interface. At the top, there is a green header bar with various icons and a user profile 'hue'. Below the header, the title 'Job Browser' is displayed. The interface includes a search bar with 'Username: hue' and 'Text: Search for text'. There is a checkbox for 'Show retired jobs' and a status filter bar with 'Succeeded', 'Running', 'Failed', and 'Killed'. The main table lists jobs with columns: Logs, ID, Name, Status, User, Maps, Reduces, Queue, Priority, Duration, and Date. The first entry is a job with ID '1425317063600\_0001', Name 'wordcount', Status 'SUCCEEDED', User 'hue', Maps '100%', Reduces '100%', Queue 'default', Priority 'N/A', Duration '6m:27s', and Date '03/02/15 09:40:20'. At the bottom, it says 'Showing 1 to 1 of 1 entries' and has navigation buttons 'Previous', '1', and 'Next'.

## Description:

Implementing a Hadoop program that takes text files as input and finds the word with maximum frequency in the whole data-set using MapReduce functionality of Hadoop.

Map Reduce has 3 phases: Map, Shuffle and Reduce.

Shuffle is inbuilt, After mapping the Intermediate results are shuffled by Hadoop.

The Mapper of Map phase takes a collection of data and convert it into another set of data, where individual elements are broken into <key,value> pairs.

These <key,value> pairs are shuffled by shuffle function and passed to Reducer.

The Reducer of Reduce phase takes the output from map phase as input and intermediate <key,value> pairs are aggregated into smaller set of <key,value> pairs as an output.

## Map Function:

```
public void map(LongWritable key, Text value, Context context)
throws IOException, InterruptedException
```

key - input key.

value - input value.

context – write method in context is used to write the output to reducer.

*Or Instead of Context, You can use:*

*output - collects mapped keys and values. OutputCollector<key,value> adds a key/value pair to the output.*

*reporter - To report progress or just indicate that they are alive*

```
String line = value.toString();
StringTokenizer tokenizer = new StringTokenizer(line);
```

Converting value to string and storing it in a line variable and passing as an argument to tokenizer.

```
while (tokenizer.hasMoreTokens())
{
    word.set(tokenizer.nextToken());
    context.write(word, one);
}
```

If more tokens available, then enters while loop. In while loop, we are passing next available token to word, which is a new value object.

`context.write(word, one)` collects <key,value> pair and increments if it encounters that word in the file and sends as output.

For example:

Input: hi hi hi

```
context.write(hi, one) = hi [1,1,1]
```

Because we set one= IntWritable[1]

## Reduce Function:

```
public void reduce(Text key, Iterable<IntWritable> values, Context context)
throws IOException, InterruptedException
```

key - Input key from map.

Iterator – For Iterating the count of one word found

values - input value from map.

context – write method in context is used to write the output.

*Or Instead of Context, You can use:*

*output - collects mapped keys and values. OutputCollector<key,value> adds a key/value pair to the output.*

*reporter - To report progress or just indicate that they are alive*

```
for(IntWritable value: values)
{
    sum+= value.get();
}
```

For each value in values, enters for each() loop and increments the sum for each occurrence of word in file. i.e. how many times each word is appeared.

```
if (sum > max_sum)
{
    max_sum = sum;
    max_occured_key.set(key);
}
```

To find the word with highest occurrence, Initializing max\_sum to 0 and creating new Text() variable called max\_occured\_key at the beginning. If the sum is greater than max\_sum, then assign that sum to max\_sum and set the key value as max\_occured\_key. So at the end, max\_sum will have the sum of highest occurring word in the file and max\_occured\_key will have the key value which has highest sum i.e. key with highest occurrence.

```
protected void cleanup(Context context)
throws IOException, InterruptedException
{
    context.write(max_occured_key, new IntWritable(max_sum));
}
```

cleanup(Context context) is used to print single value as output, i.e. Usually reducer prints output of all keys and values. But we are supposed to get only key with highest occurrence as output. So used cleanup.

context.write(max\_occured\_key, new IntWritable(max\_sum)) prints the key that has high occurrence and its value i.e. max\_sum.

Example: Input = hi [1,1,1],

Output = hi [3]