


## CS540 - Practice problems on indexing and query processing


1- Consider the B+ tree shown below.

- Delete keys 13 and 17 from this B+ tree.
- Insert keys 49, 51, 55 and 60 into this B+ tree.

You only need to show the final picture of the B+ tree.



**Solution:**



2- Consider the following relational schema. Primary keys are underlined and foreign keys are in italics.

Student (Student\_Id, Name)

Health\_Insurance (Insurance\_Id, Date\_Started, *Student\_Id*)

Assume that number of blocks in relation student is 10,000 and the number of blocks in relation Health\_Insurance is 300. We do not have any indexes on these relations. We like to join these relations on Student\_ID.

- a) If we can have 12,000 blocks in main memory, i.e.,  $M = 20,000$ , explain what is the fastest join algorithm to join “Student” and “Health\_Insurance” on Student\_Id and analyze its cost?

Solution: Because we can fit both relations in main memory, we may use an internal memory join algorithms to join them. We may pick the internal memory version of nested-loop, sort-merge, or hash-join as they involve equal number of I/O accesses. The cost of join will be  **$B(\text{Health\_Insurance}) + B(\text{Student})$**  disk I/O's. Note that we do not care about the cost of memory operations when analyzing the cost of query processing algorithms.

- b) If we can fit 20 blocks in main memory,  $M = 20$ , what is the fastest join algorithm to join these relations? Analyze its cost.

Solution: Hash-join and optimized sort-merge join are the fastest algorithms in the given setting and have equal costs. However, to apply optimized sort-merge join, we should have  **$B(\text{Health\_Insurance}) + B(\text{Student}) \leq M^2$** . As  **$B(\text{Health\_Insurance}) + B(\text{Student})$**  is not less than or equal to 400, we cannot use optimized sort-merge algorithm. Hash-join

algorithm, however, requires smaller amount of main memory. Because the smaller relation is **Health\_Insurance**, we need to have  $B(\text{Health\_Insurance}) \leq M^2$  or  $B(\text{Health\_Insurance}) \leq 400$  to perform this join using hash-join algorithm. As we have  $B(\text{Health\_Insurance}) \leq 400$ , we can use hash-join algorithm for the join. The cost of hash-join for this join is  $3 B(\text{Health\_Insurance}) + 3 B(\text{Student}) = 30900$  number of I/O access.

**3-** Assume that we like to sort a relation R. Answer the following questions.

- a) If R is too large to fit in the main memory and does not have any index, explain which algorithm we can use to sort R and analyze its cost for the sort.

**Solution:**

We may use two-pass, multiway merge-sort algorithm. The cost of sort is  $3B(R)$ .

- b) Assuming we can fit 200 blocks in main memory,  $M= 200$ , what should be the size of relation R (in blocks) to use the algorithm of part a?

**Solution:**

A relation of B blocks can be sorted using the algorithms of part a, as long as B is no more than  $M^2$ . Hence, if the size of R is less than or equal to 40,000, we can use two-phase, multiway merge-sort algorithm to sort it.